

USASK Data Science Bootcamp T6: Statistical Machine Learning

Ridge, Lasso, Elastic Net R Tutorial

TA: Kyle Gardiner Tutors: Lina Li and Jing Wang

June 13, 2023. (3:50pm-4:30pm)

This document will outline how to employ Ridge regression, Lasso Regression, and Elastic Net. First we will load in the necessary packages for this tutorial.

```
rm(list=ls(all=TRUE)) # removes objects from the environment to start fresh

library(ISLR2) # contains the dataset "Auto" used in this tutorial
library(glmnet) # used for ridge, lasso, and elastic net
library(caret) # used for elastic net
```

We will be using the same dataset for all 3 methods so lets load in the data and split it before going into each of the methods.

Loading the data

```
data(Auto) # adds the "Auto" dataset to the environment

?Auto # Prints dataframe information in the 'help' console

str(Auto) # Displays the structure of the "Auto" dataset
```

```
## 'data.frame':   392 obs. of  9 variables:
##  $ mpg          : num  18 15 18 16 17 15 14 14 14 15 ...
##  $ cylinders    : int   8  8  8  8  8  8  8  8  8  8 ...
##  $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
##  $ horsepower  : int  130 165 150 150 140 198 220 215 225 190 ...
##  $ weight       : int 3504 3693 3436 3433 3449 4341 4354 4312 4425 3850 ...
##  $ acceleration: num   12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
##  $ year         : int   70 70 70 70 70 70 70 70 70 70 ...
##  $ origin       : int    1  1  1  1  1  1  1  1  1  1 ...
##  $ name         : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231 14 161 141 54 223 241 ...
##  - attr(*, "na.action")= 'omit' Named int [1:5] 33 127 331 337 355
##  ..- attr(*, "names")= chr [1:5] "33" "127" "331" "337" ...
```

```
dim(Auto) # Displays the dimensions of the "Auto" dataset (rows x columns)
```

```
## [1] 392 9
```

```
sum(is.na(Auto)) # checking for missing values that might need to be dealt with before data analysis
```

```
## [1] 0
```

Formatting data

Before we start building the model, we need to do some data formatting that will help us in the later sections.

```
# need to convert categorical variables to factors
Auto$cylinders<-as.factor(Auto$cylinders)
Auto$year<-as.factor(Auto$year)
Auto$origin<-as.factor(Auto$origin)

# using 'model.matrix()' to create a design matrix
xx <- model.matrix(mpg ~., Auto[, -9]) # using all the predictors except for 'name'
yy <- Auto$mpg # defining the outcome variable; 'mpg'

# specifying grid to be used later in the functions
grid<-10^seq(10, -2, length=100)
```

Splitting the Data

```
set.seed(1) # set the seed for reproducibility
id<-sample(nrow(Auto), round(nrow(Auto)*0.80)) # Used to identify indexes that fall in train or test
xx.train<-xx[id,] # 'xx.train' contains 80% of the original data for predictors
xx.test<-xx[-id,] # 'xx.test' contains the remaining 20% of the original data for predictors
yy.train<-yy[id] # 'yy.train' contains 80% of the original data for outcome
yy.test<-yy[-id] # 'yy.test' contains the remaining 20% of the original data for outcome
```


Once we have finished loading, formatting, and splitting the data, we can start with ridge regression.

Ridge Regression

(a) Ridge Regression Model Building

```
# To fit the ridge regression model, we use the 'glmnet()' function from the glmnet package
ridge.mod <- glmnet(xx.train, yy.train, alpha = 0, lambda = grid)
# IMPORTANT: specifying 'alpha' is necessary to apply the proper regularization penalty to the model
# 'alpha=0' specifies ridge regression

plot(ridge.mod, label = T)
```




plot to visualize how the coefficient estimates change as the shrinkage penalty changes

(b) Cross Validation to identify best λ that reduces error

```
ridge.cv <- cv.glmnet(xx.train, yy.train, alpha = 0, lambda = grid, nfolds=10)
# 10 fold cross validation

plot(ridge.cv) # plot to visualize how MSE changes as lambda changes
```



(c) Assessing model accuracy

```
bestlam.ridge <- ridge.cv$lambda.1se # can either use 'lambda.min' or 'lambda.1se'

ridge.pred=predict(ridge.mod, s=bestlam.ridge, newx=xx.test)
# predicting the values using the non-cv ridge model and the best lambda from cross validation

mean((ridge.pred-yy.test)^2) # calculating mean squared error for this model
```

```
## [1] 7.698884
```

We will compare the accuracy of these models at the very end.

Lasso Regression

(a) Lasso Regression Model Building

```
# To fit the lasso regression model, we use the 'glmnet()' function from the glmnet package
lasso.mod <- glmnet(xx.train, yy.train, alpha = 1, lambda = grid)
# IMPORTANT: specifying 'alpha' is necessary to apply the proper regularization penalty to the model
```

```


# 'alpha=1' specifies lasso regression
plot(lasso.mod, label = T)

```

```

## Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
## collapsing to unique 'x' values

```



```

# plot to visualize how the coefficient estimates change as the shrinkage penalty changes

```


(b) Cross Validation to identify best λ that reduces error

```

lasso.cv <- cv.glmnet(xx.train, yy.train, alpha = 1, lambda = grid, nfolds=10)
# 10 fold cross validation

plot(lasso.cv) # plot to visualize how MSE changes as lambda changes

```



(c) Assessing model accuracy

```
bestlam.lasso <- lasso.cv$lambda.1se # can either use 'lambda.min' or 'lambda.1se'

lasso.pred=predict(lasso.mod, s=bestlam.lasso, newx=xx.test)
# predicting the values using the non-cv lasso model and the best lambda from cross validation

mean((lasso.pred-yy.test)^2) # calculating mean squared error for this model

## [1] 7.476549
```

Again, we will compare the performance of all the model at the end.

Elastic Net

Elastic Net combines Ridge and Lasso regression.

(a) Cross Validation

In this section, we use cross validation to find the most appropriate model for our data. Cross validation allows us to select the best α and λ to reduce the error in the model

```
# trainControl allows us to define control parameters to use in the following 'train()' function
cv_10 = trainControl(method = "cv", number = 10) # here we specify 10 fold cross validation

set.seed(1) # set seed for reproducibility

# Using the 'train()' function, we train a cross validation elastic net
elastic.net.cv = train(
  mpg ~., # formula to predict mpg
  data = Auto[id,-9], # using the training data with all the predictor, except for 'name'
  method = "glmnet", # 'method="glmnet"' specifies using elastic net
  trControl = cv_10) # using the control parameters we specified earlier
```

(b) Determining Best α and λ

After fitting the cross validation elastic net, we can extract the values of α and λ that correspond to the smallest Root Mean Squared Error (RMSE).

```
elastic.net.cv # prints table for each combination of alpha and lambda and corresponding RMSE
```

```
## glmnet
##
## 314 samples
## 7 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 284, 283, 283, 281, 283, 282, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda      RMSE      Rsquared  MAE
##  0.10   0.01294453  3.042737  0.8524046  2.308646
##  0.10   0.12944527  3.067450  0.8506485  2.313612
##  0.10   1.29445274  3.234601  0.8394891  2.440086
##  0.55   0.01294453  3.046206  0.8520572  2.307647
##  0.55   0.12944527  3.087948  0.8493096  2.320911
##  0.55   1.29445274  3.704561  0.8051795  2.803092
##  1.00   0.01294453  3.050636  0.8517215  2.308897
##  1.00   0.12944527  3.113617  0.8476793  2.338149
##  1.00   1.29445274  4.163707  0.7677863  3.142528
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 0.1 and lambda = 0.01294453.
```

```
# extracting optimal alpha and lambda
myid <- order(elastic.net.cv$results$RMSE)[1] # determines the index that has the smallest RMSE
```

```
# extracts the alpha that corresponds to the smallest RMSE
mybest.alpha <- elastic.net.cv$results$alpha[myid]

# extracts the lambda that corresponds to the smallest RMSE
mybest.lambda <- elastic.net.cv$results$lambda[myid]
```

(c) Fitting Elastic Net with optimal α and λ

```
# to do this, use the 'glmnet()' function
elastic.net.model<- glmnet(xx.train, yy.train, alpha = mybest.alpha, lambda = mybest.lambda)
# using 'mybest.alpha' and 'mybest.lambda' from cross validation

coef(elastic.net.model)
```

```
## 24 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept) 30.805500302
## (Intercept) .
## cylinders4  6.247375595
## cylinders5  5.235532941
## cylinders6  3.995908710
## cylinders8  6.102187506
## displacement 0.007275972
## horsepower  -0.038544698
## weight      -0.004997299
## acceleration 0.053986586
## year71       1.255078477
## year72      -0.294989347
## year73      -0.412742892
## year74       1.276020491
## year75       0.816982936
## year76       0.994356936
## year77       3.200435545
## year78       3.281689080
## year79       4.767955393
## year80       9.788016445
## year81       6.630889068
## year82       8.070358153
## origin2      1.978721333
## origin3      1.885491137
```

```
# prints the coefficient estimates for the variables used in the elastic net model
```


(d) **Assessing model accuracy**

Once we have built the elastic net model, we can assess its accuracy.

```
# Again, we predict the test outcomes using the built elastic net model
elastic.net.pred<-predict(elastic.net.model, s=mybest.lambda, newx=xx.test)
# need to specify 's=mybest.lambda' to ensure we are using the proper lambda to predict

# Now we can find the Mean Squared Error for this model
mean((elastic.net.pred-yy.test)^2)
```

```
## [1] 7.656109
```

Final Comparisons

Lasso MSE: 7.698884

Ridge MSE: 7.476549

Elastic Net MSE: 7.656109

This concludes the tutorial of how to employ ridge regression, lasso regression, and elastic net in R.